

# Cornell Bowers C-IS


College of Computing  
and Information Science

## Toward More Robust Hyperparameter Optimization

**A. Feder Cooper**

Ph.D. Student, Department of Computer Science

# Brief bio



JZ Forde\*, **AF Cooper**\*, et. al. [SEDL@ICLR '21]


**AF Cooper** and E Abrams [AIES '21]

**AF Cooper**, et. al. [NeurIPS '21]




R Zhang, **AF Cooper**, et. al. [AISTATS '20]

R Zhang\*, **AF Cooper**\*, et. al. [NeurIPS '20]



# Brief bio



JZ Forde\*, **AF Cooper**\*, et. al. [SEDL@ICLR '21]


**AF Cooper** and E Abrams [AIES '21]

**AF Cooper**, et. al. [NeurIPS '21]



R Zhang, **AF Cooper**, et. al. [AISTATS '20]

R Zhang\*, **AF Cooper**\*, et. al. [NeurIPS '20]



# Hyperparameter Optimization Is Deceiving Us, and How to Stop It



AF Cooper, et. al. [NeurIPS '21]

Joint work with



Yucheng Lu, Jessica Zosa Forde, and Chris De Sa



# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform




# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform




non-adaptive optimizers  
outperform adaptive optimizers

# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform




non-adaptive optimizers  
outperform adaptive optimizers




non-adaptive optimizers  
do not outperform adaptive optimizers

# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform



non-adaptive optimizers  
outperform adaptive optimizers


$p$



non-adaptive optimizers  
do not outperform adaptive optimizers

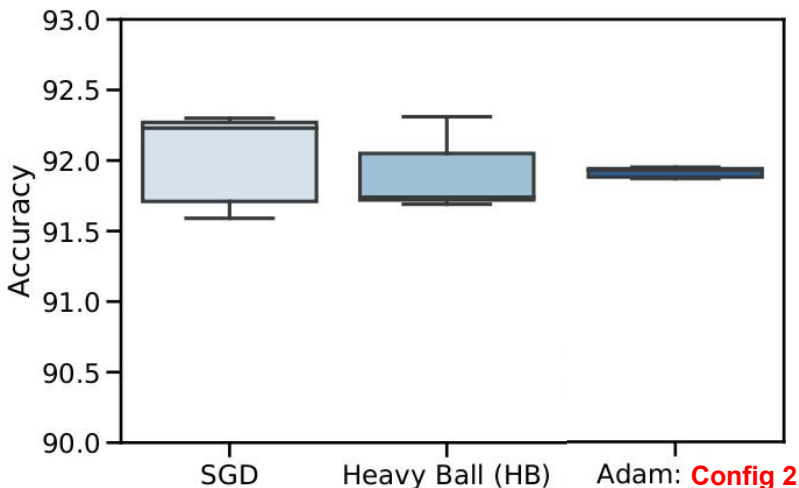


# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform



non-adaptive optimizers  
outperform adaptive optimizers


$p$



non-adaptive optimizers  
do not outperform adaptive optimizers

$\neg p$


# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform



non-adaptive optimizers  
outperform adaptive optimizers

$p$

AND





non-adaptive optimizers  
do not outperform adaptive optimizers

$\neg p$

→ FALSE



# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform

We do not know the ground truth



# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform



We do not know the ground truth



It is **fine** for the ML community to accept **either** to form **conclusions**

# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform

We do not know the ground truth





It is **fine** for the ML community to accept **either** to form **conclusions**

It is **fine** for the ML community to accept **neither** to form **no conclusions**

# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform

We do not know the ground truth




It is **fine** for the ML community to accept **either** to form **conclusions**

It is **fine** for the ML community to accept **neither** to form **no conclusions**

It is **not fine** for the ML community to accept **both** to form **inconsistent conclusions**

# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform



We do not know the ground truth



$p$

# We want to prevent our conclusions from depending on the underlying configuration of the HPO we perform

We do not know the ground truth



We do not want this to be possible

$p$

$\neg p$



# We call this phenomenon “hyperparameter deception”

No one has studied HPO at this meta level of “possible” outcomes before

For **ML practitioners**, the **goal is to deploy a specific model**  
pick a few hyperparameters to test, compare results, and  
deploy the model that achieved the best performance

For **ML research**, the **goal is to derive general knowledge**  
pick a few hyperparameters to test, compare results, and  
deploy the model that achieved the best performance

# We call this phenomenon “hyperparameter deception”

No one has studied HPO at this meta level of “possible” outcomes before

For ML practitioners, the goal is to deploy a specific model  
pick a few hyperparameters to test, compare results, and  
deploy the model that achieved the best performance

Doing the  
same thing to  
achieve very  
different goals

For ML research, the goal is to derive general knowledge  
pick a few hyperparameters to test, compare results, and  
deploy the model that achieved the best performance

The process of picking HPO configurations to test is vague

Whether we believe our conclusions or not is also vague

# We formally study “hyperparameter deception”

The paper formalizes the problem in 2 phases:

- 1) We remove the vagueness from the problem by pinning down
  - a) a concrete formalization for the **possible outcomes of running hyperparameter optimization**
  - b) a concrete formalization of our **belief in conclusions**
- 2) Proving non-trivial theorems about whether a hyperparameter optimization procedure is defended against “deception” (and validating this empirically)

# We formally study “hyperparameter deception”

The paper formalizes the problem in 2 phases:

- 1) **We remove the vagueness from the problem by pinning down**
  - a) a concrete formalization for the **possible outcomes of running hyperparameter optimization**
  - b) a concrete formalization of our **belief in conclusions**
- 2) Proving non-trivial theorems about whether a hyperparameter optimization procedure is defended against “deception” (and validating this empirically)

# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ .

# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.

# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.



# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.



# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.

# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.

# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.



# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.

# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.

# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.



# We formalize the process of drawing conclusions from empirical studies using HPO

We formalize all of the randomness in HPO in terms of a random seed  $r$  and a pseudo-random number generator (PRNG)  $G$ . Given a seed,  $G$  deterministically produces a sequence of pseudo-random numbers: all numbers lie in some set  $\mathcal{I}$  (typically 64-bit integers), i.e.  $r \in \mathcal{I}$  and PRNG  $G : \mathcal{I} \rightarrow \mathcal{I}^\infty$ . With this, we can now define HPO formally:

**Definition 2.** An HPO procedure  $H$  is a tuple  $(H_*, \mathcal{C}, \Lambda, \mathcal{A}, \mathcal{M}, G, X)$  where  $H_*$  is a randomized algorithm,  $\mathcal{C}$  is a set of allowable hyper-HPs (i.e., allowable configurations for  $H_*$ ),  $\Lambda$  is a set of allowable HPs (i.e., of HP sets  $\lambda$ ),  $\mathcal{A}$  is a training algorithm (e.g. SGD),  $\mathcal{M}$  is a model (e.g. VGG16),  $G$  is a PRNG, and  $X$  is some dataset (usually split into train and validation sets). When run,  $H_*$  takes as input a hyper-HP configuration  $c \in \mathcal{C}$  and a random seed  $r \in \mathcal{I}$ , then proceeds to run  $\mathcal{A}_\lambda$  (on  $\mathcal{M}_\lambda$  using  $G(r)$  and data<sup>3</sup> from  $X$ ) some number of times for different HPs  $\lambda \in \Lambda$ . Finally,  $H_*$  outputs a tuple  $(\lambda^*, \ell)$ , where  $\lambda^*$  is the HP configuration chosen by HPO and  $\ell$  is the log documenting the run.

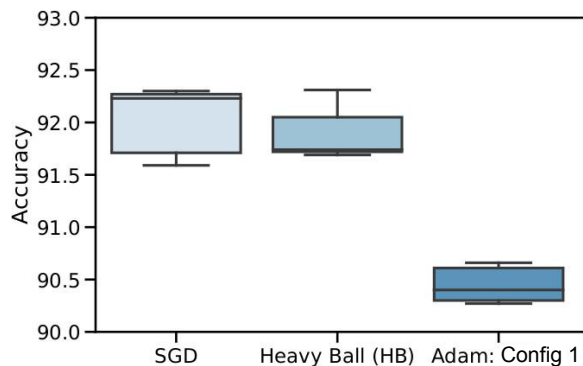
# We formalize the process of drawing conclusions from empirical studies using HPO

**Definition 1.** *A log  $\ell$  records all the choices and measurements made during an HPO run, including the total time  $T$  it took to run. It has all necessary information to make the HPO run reproducible.*



# We formalize the process of drawing conclusions from empirical studies using HPO


**Definition 1.** A log  $\ell$  records all the choices and measurements made during an HPO run, including the total time  $T$  it took to run. It has all necessary information to make the HPO run reproducible.



3 HPO procedures;  
Each box plot corresponds to a log

# We formalize the process of drawing conclusions from empirical studies using HPO

**Definition 1.** A log  $\ell$  records all the choices and measurements made during an HPO run, including the total time  $T$  it took to run. It has all necessary information to make the HPO run reproducible.



3 HPO procedures;  
Each box plot corresponds to a **log**

# We then formalize the process of drawing conclusions from empirical studies using HPO

**Definition 3.** An *epistemic hyperparameter optimization procedure (EHPO)* is a tuple  $(\mathcal{H}, \mathcal{F})$  where  $\mathcal{H}$  is a set of HPO procedures  $H$  (Definition 2) and  $\mathcal{F}$  is a function that maps a set of HPO logs  $\mathcal{L}$  (Definition 1) to a set of logical formulas  $\mathcal{P}$ , i.e.  $\mathcal{F}(\mathcal{L}) = \mathcal{P}$ . An execution of EHPO involves running each  $H \in \mathcal{H}$  some number of times (each run produces a log  $\ell$ ), and then evaluating  $\mathcal{F}$  on the logs  $\mathcal{L}$  produced in order to output the conclusions  $\mathcal{F}(\mathcal{L})$  we draw from all of the HPO runs.

# We then formalize the process of drawing conclusions from empirical studies using HPO


**Definition 3.** An *epistemic hyperparameter optimization procedure (EHPO)* is a tuple  $(\mathcal{H}, \mathcal{F})$  where  $\mathcal{H}$  is a set of HPO procedures  $H$  (Definition 2) and  $\mathcal{F}$  is a function that maps a set of HPO logs  $\mathcal{L}$  (Definition 1) to a set of logical formulas  $\mathcal{P}$ , i.e.  $\mathcal{F}(\mathcal{L}) = \mathcal{P}$ . An execution of EHPO involves running each  $H \in \mathcal{H}$  some number of times (each run produces a log  $\ell$ ), and then evaluating  $\mathcal{F}$  on the logs  $\mathcal{L}$  produced in order to output the conclusions  $\mathcal{F}(\mathcal{L})$  we draw from all of the HPO runs.

**Epistemic Hyperparameter Optimization (EHPO)** takes a **set of HPO procedures** and a function **F**, which maps a set of HPO procedure **logs** to **conclusions** about algorithm performance

# We then formalize the process of drawing conclusions from empirical studies using HPO


**Epistemic Hyperparameter Optimization (EHPO)** takes a set of HPO procedures and a function  $F$ , which maps a set of HPO procedure logs to conclusions about algorithm performance

Log set **L1**  
( $|L1| = 3$ )



$F(L1) \rightarrow$  “Non-adaptive optimizers outperform adaptive ones” =  $p$

Log set **L2**  
( $|L2| = 3$ )



$F(L2) \rightarrow$  “Non-adaptive optimizers do not outperform adaptive ones” =  $\neg p$

# We frame drawing inconsistent conclusions from EHPO in terms of an adversary who can deceive us



“I will suppose...an evil genius, supremely powerful and clever, **who has directed his entire effort at deceiving me.** I will regard the heavens, the air, the earth, colors, shapes, sounds, and all external things as nothing but the bedeviling hoaxes of my dreams, with which he lays snares for my credulity...**even if it is not within my power to know anything true, it certainly is within my power to take care resolutely to withhold my assent to what is false,** lest this deceiver, however powerful, however clever he may be, have any effect on me.”

# We frame drawing inconsistent conclusions from EHPO in terms of an adversary who can deceive us



“I will suppose...an evil genius, supremely powerful and clever, who has directed his entire effort at deceiving me. I will regard the heavens, the air, the earth, colors, shapes, sounds, and all external things as nothing but the bedeviling hoaxes of my dreams, with which he lays snares for my credulity...even if it is not within my power to know anything true, it certainly is within my power to take care resolutely to withhold my assent to what is false, lest this deceiver, however powerful, however clever he may be, have any effect on me.”

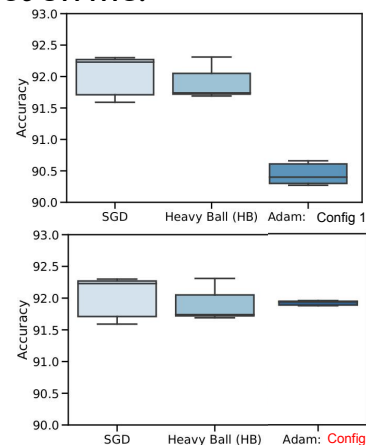
Recall:

**We do not know the ground truth**

It is **fine** to accept **either**

It is **fine** to accept **neither** (form **no conclusions**)

It is **not fine** to accept **both** (inconsistent conclusions are false)



# We frame drawing inconsistent conclusions from EHPO in terms of an adversary who can deceive us



Imagine an **evil demon** who  
is trying to deceive us about relative algorithm performance via EHPO  
maintains a set of log HPO logs which it can modify  
presents us with a final log set, from which we can draw conclusions



# We frame drawing inconsistent conclusions from EHPO in terms of an adversary who can deceive us



Imagine an **evil demon** who  
is trying to deceive us about relative algorithm performance via EHPO  
maintains a set of log HPO logs which it can modify  
presents us with a final log set, from which we can draw conclusions


**Definition 4.** A randomized **strategy**  $\sigma$  is a function that specifies which action the demon will take. Given  $\mathcal{L}$ , its current set of logs,  $\sigma(\mathcal{L})$  gives a distribution over concrete actions, where each action is either 1) running a new  $H$  with its choice of hyper-HPs  $c$  and seed  $r$  2) erasing some logs, or 3) returning. We let  $\Sigma$  denote the set of all such strategies.

# We frame drawing inconsistent conclusions from EHPO in terms of an adversary who can deceive us




Imagine an **evil demon** who  
is trying to deceive us about relative algorithm performance via EHPO  
maintains a set of HPO logs which it can modify  
presents us with a final log set, from which we can draw conclusions

*could* produce **L1**



$$F(\mathbf{L1}) \rightarrow p$$

or *could* produce **L2**



$$F(\mathbf{L2}) \rightarrow \neg p$$

# We frame drawing inconsistent conclusions from EHPO in terms of an adversary who can deceive us


Imagine an **evil demon** who

is trying to deceive us about relative algorithm performance via EHPO

maintains a set of HPO logs which it can modify

presents us with a final log set, from which we can draw conclusions

*could* produce **L1**



$$F(\mathbf{L1}) \rightarrow p$$

*or could*



$$\neg p$$

# Modal logic is the standard way to formalize “could”

**Modal logic** extends propositional logic to allow us to reason about **possibility** introducing an additional operator

$$\phi := P \mid \neg\phi \mid \phi \wedge \phi \mid \Diamond\phi$$

# Modal logic is the standard way to formalize “could”

**Modal logic** extends propositional logic to allow us to reason about **possibility** introducing an additional operator

$$\phi := P \mid \neg\phi \mid \phi \wedge \phi \mid \Diamond\phi$$

$\Diamond\phi$  reads, “It is **possible** that  $\phi$ .”

# Modal logic is the standard way to formalize “could”

**Modal logic** extends propositional logic to allow us to reason about **possibility** introducing an additional operator

$$\phi := P \mid \neg\phi \mid \phi \wedge \phi \mid \Diamond\phi$$

$\Diamond\phi$  eads, “It is **possible** that  $\phi$ .”

**We combine two modal operators** so that we can capture the idea that **it is not possible** to adopt inconsistent **beliefs** about experimental outcomes

# Modal logic is the standard way to formalize “could”

**Modal logic** extends propositional logic to allow us to reason about **possibility** introducing an additional operator

$$\phi := P \mid \neg\phi \mid \phi \wedge \phi \mid \Diamond\phi$$

$\Diamond\phi$  eads, “It is **possible** that  $\phi$ .”

**We combine two modal operators** so that we can capture the idea that it is **not possible** to adopt inconsistent **beliefs** about experimental outcomes

$$\psi := P \mid \neg\psi \mid \psi \wedge \psi \mid \Diamond_t\psi \mid \mathcal{B}\psi$$

# Expressing the possible outcomes of EHPO with $\Diamond_{tp}$

Syntax

$\Diamond_{tp}$


Intuition

The demon could adopt a strategy for running EHPO that is guaranteed to cause their desired outcome  $p$  in at most time  $t$  in expectation

Semantics

A set of EHPO output logs **models** that it is possible  $p$  time  $t$

$\Diamond_{tp}$





# Expressing the possible outcomes of EHPO with $\Diamond_t p$

## Syntax

$$\Diamond_t p$$

**Definition.** A randomized *strategy*  $\sigma$  is a function that specifies which action the demon will take. Given  $\mathcal{L}$ , its current set of logs,  $\sigma(\mathcal{L})$  gives a distribution over concrete actions, where each action is either 1) running a new  $H$  with its choice of hyper-HPs  $c$  and seed  $r$  2) erasing some logs, or 3) returning. We let  $\Sigma$  denote the set of all such strategies.

# Expressing the possible outcomes of EHPO with $\Diamond_t p$


## Syntax

$$\Diamond_t p$$


**Definition.** A randomized *strategy*  $\sigma$  is a function that specifies which action the demon will take. Given  $\mathcal{L}$ , its current set of logs,  $\sigma(\mathcal{L})$  gives a distribution over concrete actions, where each action is either 1) running a new  $H$  with its choice of hyper-HPs  $c$  and seed  $r$  2) erasing some logs, or 3) returning. We let  $\Sigma$  denote the set of all such strategies.

We can now define what the demon can reliably bring about, in terms of executing a strategy in bounded time:

# Expressing the possible outcomes of EHPO with



## Syntax



**Definition.** A randomized **strategy**  $\sigma$  is a function that specifies which action the demon will take. Given  $\mathcal{L}$ , its current set of logs,  $\sigma(\mathcal{L})$  gives a distribution over concrete actions, where each action is either 1) running a new  $H$  with its choice of hyper-HPs  $c$  and seed  $r$  2) erasing some logs, or 3) returning. We let  $\Sigma$  denote the set of all such strategies.

We can now define what the demon can reliably bring about, in terms of executing a strategy in bounded time:

**Definition.** Let  $\sigma[\mathcal{L}]$  denote the logs output from executing strategy  $\sigma$  on logs  $\mathcal{L}$ , and let  $\tau_\sigma(\mathcal{L})$  denote the total time spent during execution.  $\tau_\sigma(\mathcal{L})$  is equivalent to the sum of the times  $T$  it took each HPO procedure  $H \in \mathcal{H}$  executed in strategy  $\sigma$  to run. Note that both  $\sigma[\mathcal{L}]$  and  $\tau_\sigma(\mathcal{L})$  are random variables, as a function of the randomness of selecting  $G$  and the actions sampled from  $\sigma(\mathcal{L})$ . For any formula  $p$  and any  $t \in \mathbb{R}_{>0}$ , we say  $\mathcal{L} \models \Diamond_t p$ , i.e. “ $\mathcal{L}$  models that it is possible  $p$  in time  $t$ ,” if

there exists a strategy  $\sigma \in \Sigma$ , such that  $\mathbb{P}(\sigma[\mathcal{L}] \models p) = 1$  and  $\mathbb{E}[\tau_\sigma(\mathcal{L})] \leq t$ .

# Expressing the possible outcomes of EHPO with $\Diamond_{tp}$

Syntax

$\Diamond_{tp}$


Intuition

The demon could adopt a strategy for running EHPO that is guaranteed to cause their desired outcome  $p$  in at most time  $t$  in expectation

Semantics

A set of EHPO output logs **models** that it is possible  $p$  time  $t$

$\Diamond_{tp}$



# Expressing our belief with $Bp$

Syntax

$Bp$

Intuition

We believe/ conclude  $p$

Semantics

The set of EHPO output logs models our belief in  $p$

# Expressing our belief with $\mathcal{B}p$

Syntax

$\mathcal{B}p$

Intuition

We believe/ conclude  $p$

Semantics

The set of EHPO output logs models our belief in  $p$

**Definition 6.** For any formula  $p$ , we say  $\mathcal{L} \models \mathcal{B}p$ , “ $\mathcal{L}$  models our belief in  $p$ ”, if  $p \in \mathcal{F}(\mathcal{L})$ .

# With both of these operators, we can formalize the problem of *hyperparameter deception*

*t*-non-deceptive axiom

$$\neg (\Diamond_t \mathcal{B}p \wedge \Diamond_t \mathcal{B}\neg p)$$



If it is **possible** for the demon can get us to **believe**  $p$  in time  $t$ , then it is **not possible** for the demon to get us to **believe**  $\neg p$  in time  $t$

# With both of these operators, we can formalize the problem of *hyperparameter deception*

**t-non-deceptive axiom**

$$\neg (\Diamond_t \mathcal{B}p \wedge \Diamond_t \mathcal{B}\neg p)$$

If it is **possible** for the **demon** can get us to **believe**  $p$  in time  $t$ , then it is **not possible** for the **demon** to get us to **believe**  $\neg p$  in time  $t$



Our motivating example is **t-deceptive**



# We formally study “hyperparameter deception”

The paper formalizes the problem in 2 phases:


- 1) We remove the vagueness from the problem by pinning down
  - a) a concrete formalization for the possible outcomes of running hyperparameter optimization
  - b) a concrete formalization of our belief in conclusions
- 2) Proving non-trivial theorems about whether a hyperparameter optimization procedure is defended against “deception” (and validating this empirically)

# Using this formalization, we can prove $t$ -non-deceptive EHPO

We can  
suggest concrete EHPO

prove that this EHPO satisfies  $t$ -non-deceptiveness

Which means, for this EHPO, in time  $t$  deception is not possible



# Using this formalization, we can prove $t$ -non-deceptive EHPO

## Defense intuition:

Given some **naive reasoner**, we construct a **defended reasoner** that is always **more skeptical** than the naive reasoner

If the naive reasoner is  $t$ -non-deceptive, then any more skeptical reasoner is also  $t$ -non-deceptive

# Using this formalization, we can prove $t$ -non-deceptive EHPO

construct  $\mathcal{B}_*$  such that for any  $p$ ,  $\mathcal{B}_*p \equiv \mathcal{B}_n p \wedge \neg \Diamond_t \mathcal{B}_n \neg p$  (1).

Directly from our axioms (Section 4), we can now prove  $\mathcal{B}_*$  is defended. We will suppose it is possible for  $\mathcal{B}_*$  to be deceived, demonstrate a contradiction, and thereby guarantee that  $\mathcal{B}_*$  is  $t$ -non-deceptive. Suppose  $\mathcal{B}_*$  can be deceived in time  $t$ , i.e.  $\Diamond_t \mathcal{B}_* p \wedge \Diamond_t \mathcal{B}_* \neg p$  is True. Starting with the left,  $\Diamond_t \mathcal{B}_* p$ :

Rule	
$\Diamond_t \mathcal{B}_* p \equiv \Diamond_t (\mathcal{B}_n p \wedge \neg \Diamond_t \mathcal{B}_n \neg p)$	Applying $\Diamond_t$ to the definition of $\mathcal{B}_* p$ (I)
$\rightarrow \Diamond_t (\neg \Diamond_t \mathcal{B}_n \neg p)$	Reducing a conjunction to either of its terms: $(a \wedge b) \rightarrow b$
$\rightarrow \neg \Diamond_t \mathcal{B}_n \neg p$	Symmetry; dropping all but the right-most operator: $\Diamond_t(\Diamond_t a) \rightarrow \Diamond_t a$

We then pause to apply our axioms to the right side of the conjunction,  $\Diamond_t \mathcal{B}_* \neg p$ :

Rule	
$\Diamond_t \mathcal{B}_* \neg p \equiv \Diamond_t (\mathcal{B}_n \neg p \wedge \neg \Diamond_t \mathcal{B}_n p)$	Applying $\Diamond_t$ to the definition of $\mathcal{B}_* \neg p$ (I)
$\rightarrow \Diamond_t \mathcal{B}_n \neg p \wedge \Diamond_t \neg \Diamond_t \mathcal{B}_n p$	Distributing $\Diamond_t$ over $\wedge$ : $\Diamond_t(a \wedge b) \rightarrow (\Diamond_t a \wedge \Diamond_t b)$
$\rightarrow \Diamond_t \mathcal{B}_n \neg p$	Reducing a conjunction to either of its terms: $(a \wedge b) \rightarrow a$

We now bring both sides of the conjunction back together:  $\Diamond_t \mathcal{B}_* p \wedge \Diamond_t \mathcal{B}_* \neg p \equiv \neg \Diamond_t \mathcal{B}_n \neg p \wedge \Diamond_t \mathcal{B}_n \neg p$ . The right-hand side is of the form  $\neg a \wedge a$ , which must be False. This contradicts our initial assumption that  $\mathcal{B}_*$  is  $t$ -deceptive (i.e.,  $\Diamond_t \mathcal{B}_* p \wedge \Diamond_t \mathcal{B}_* \neg p$  is True). Therefore,  $\mathcal{B}_*$  is  $t$ -non-deceptive.

# Using this formalization, we can prove $t$ -non-deceptive EHPO

construct  $\mathcal{B}_*$  such that for any  $p$ ,  $\mathcal{B}_*p \equiv \mathcal{B}_n p \wedge \neg \Diamond_t \mathcal{B}_n \neg p$  (1).

Directly from our axioms (Section 4), we can now prove  $\mathcal{B}_*$  is defended. We will suppose it is possible for  $\mathcal{B}_*$  to be deceived, demonstrate a contradiction, and thereby guarantee that  $\mathcal{B}_*$  is  $t$ -non-deceptive. Suppose  $\mathcal{B}_*$  can be deceived in time  $t$ , i.e.  $\Diamond_t \mathcal{B}_* p \wedge \Diamond_t \mathcal{B}_* \neg p$  is True. Starting with the left,  $\Diamond_t \mathcal{B}_* p$ :

**Defense takeaway:**

It is always possible to construct a  $t$ -defended EHPO

We then

Rule	
$\Diamond_t \mathcal{B}_* \neg p \equiv \Diamond_t (\mathcal{B}_n \neg p \wedge \neg \Diamond_t \mathcal{B}_n p)$	Applying $\Diamond_t$ to the definition of $\mathcal{B}_* \neg p$ (1)
$\rightarrow \Diamond_t \mathcal{B}_n \neg p \wedge \Diamond_t \neg \Diamond_t \mathcal{B}_n p$	Distributing $\Diamond_t$ over $\wedge$ : $\Diamond_t(a \wedge b) \rightarrow (\Diamond_t a \wedge \Diamond_t b)$
$\rightarrow \Diamond_t \mathcal{B}_n \neg p$	Reducing a conjunction to either of its terms: $(a \wedge b) \rightarrow a$

We now bring both sides of the conjunction back together:  $\Diamond_t \mathcal{B}_* p \wedge \Diamond_t \mathcal{B}_* \neg p \equiv \neg \Diamond_t \mathcal{B}_n \neg p \wedge \Diamond_t \mathcal{B}_n \neg p$ . The right-hand side is of the form  $\neg a \wedge a$ , which must be False. This contradicts our initial assumption that  $\mathcal{B}_*$  is  $t$ -deceptive (i.e.,  $\Diamond_t \mathcal{B}_* p \wedge \Diamond_t \mathcal{B}_* \neg p$  is True). Therefore,  $\mathcal{B}_*$  is  $t$ -non-deceptive.

# Our $t$ -defended random-search EHPO

**A defended random search EHPO.** Random search takes two hyper-HPs, a distribution  $\mu$  over the HP space and a number of trials  $K \in \mathbb{N}$  to run. HPO consists of  $K$  independent trials of training algorithms  $\mathcal{A}_{\lambda_1}, \mathcal{A}_{\lambda_2}, \dots, \mathcal{A}_{\lambda_K}$ , where the HPs  $\lambda_k$  are independently drawn from  $\mu$ , taking expected time proportional to  $K$ . When drawing conclusions, we usually look at the “best” run for each algorithm. For simplicity, we suppose there is only one algorithm,  $\mathcal{A}$ . We bound how much the choice of hyper-HPs can affect the HPs, and define a defended EHPO based on a variant of random search.

# Our $t$ -defended random-search EHPO

**A defended random search EHPO.** Random search takes two hyper-HPs, a distribution  $\mu$  over the HP space and a number of trials  $K \in \mathbb{N}$  to run. HPO consists of  $K$  independent trials of training algorithms  $\mathcal{A}_{\lambda_1}, \mathcal{A}_{\lambda_2}, \dots, \mathcal{A}_{\lambda_K}$ , where the HPs  $\lambda_k$  are independently drawn from  $\mu$ , taking expected time proportional to  $K$ . When drawing conclusions, we usually look at the “best” run for each algorithm. For simplicity, we suppose there is only one algorithm,  $\mathcal{A}$ . We bound how much the choice of hyper-HPs can affect the HPs, and define a defended EHPO based on a variant of random search.



# Our $t$ -defended random-search EHPO

**A defended random search EHPO.** Random search takes two hyper-HPs, a distribution  $\mu$  over the HP space and a number of trials  $K \in \mathbb{N}$  to run. HPO consists of  $K$  independent trials of training algorithms  $\mathcal{A}_{\lambda_1}, \mathcal{A}_{\lambda_2}, \dots, \mathcal{A}_{\lambda_K}$ , where the HPs  $\lambda_k$  are independently drawn from  $\mu$ , taking expected time proportional to  $K$ . When drawing conclusions, we usually look at the “best” run for each algorithm. For simplicity, we suppose there is only one algorithm,  $\mathcal{A}$ . We bound how much the choice of hyper-HPs can affect the HPs, and define a defended EHPO based on a variant of random search.



# Our $t$ -defended random-search EHPO

**A defended random search EHPO.** Random search takes two hyper-HPs, a distribution  $\mu$  over the HP space and a number of trials  $K \in \mathbb{N}$  to run. HPO consists of  $K$  independent trials of training algorithms  $\mathcal{A}_{\lambda_1}, \mathcal{A}_{\lambda_2}, \dots, \mathcal{A}_{\lambda_K}$ , where the HPs  $\lambda_k$  are independently drawn from  $\mu$ , taking expected time proportional to  $K$ . When drawing conclusions, we usually look at the “best” run for each algorithm. For simplicity, we suppose there is only one algorithm,  $\mathcal{A}$ . We bound how much the choice of hyper-HPs can affect the HPs, and define a defended EHPO based on a variant of random search.

# Our $t$ -defended random-search EHPO

**Definition 7.** Suppose that we are given a naive EHPO procedure  $(\{H\}, \mathcal{F}_n)$ , in which  $H$  is random search and is the only HPO in our EHPO, and  $\mathcal{F}_n$  is a “naive” belief function associated with a naive reasoner  $\mathcal{B}_n$ . For any  $K, R \in \mathbb{N}$ , we define the “ $(K, R)$ -defended” belief function  $\mathcal{F}_*$  for a skeptical reasoner  $\mathcal{B}_*$  as the following conclusion-drawing procedure. First,  $\mathcal{F}_*$  only makes conclusion set  $\mathcal{P}_*$  from a single log  $\hat{\ell}$  with  $K * R$  trials; otherwise, it concludes nothing, outputting  $\emptyset$ . Second,  $\mathcal{F}_*$  splits the single  $\hat{\ell}$  into  $R$  logs  $\ell_1, \ell_2, \dots, \ell_R$ , each containing  $K$  independent-random-search trials.<sup>8</sup> Finally,  $\mathcal{F}_*$  outputs the intersection of what the naive reasoner would have output on each log  $\ell_i$ ,

$$\mathcal{F}_*(\{\hat{\ell}\}) = \mathcal{P}_* \equiv \mathcal{F}_n(\{\ell_1\}) \cap \mathcal{F}_n(\{\ell_2\}) \cap \dots \cap \mathcal{F}_n(\{\ell_R\}).$$

Equivalently,  $\{\hat{\ell}\} \models \mathcal{B}_*p$  only if  $\{\ell_i\} \models \mathcal{B}_np$  for all  $i$ .

# Our $t$ -defended random-search EHPO

**Definition 7.** *Suppose that we are given a naive EHPO procedure  $(\{H\}, \mathcal{F}_n)$ , in which  $H$  is random search and is the only HPO in our EHPO, and  $\mathcal{F}_n$  is a “naive” belief function associated with a naive reasoner  $\mathcal{B}_n$ . For any  $K, R \in \mathbb{N}$ , we define the “ $(K, R)$ -defended” belief function  $\mathcal{F}_*$  for a skeptical reasoner  $\mathcal{B}_*$  as the following conclusion-drawing procedure. First,  $\mathcal{F}_*$  only makes conclusion set  $\mathcal{P}_*$  from a single log  $\hat{\ell}$  with  $K * R$  trials; otherwise, it concludes nothing, outputting  $\emptyset$ . Second,  $\mathcal{F}_*$  splits the single  $\hat{\ell}$  into  $R$  logs  $\ell_1, \ell_2, \dots, \ell_R$ , each containing  $K$  independent-random-search trials.<sup>8</sup> Finally,  $\mathcal{F}_*$  outputs the intersection of what the naive reasoner would have output on each log  $\ell_i$ ,*

$$\mathcal{F}_*(\{\hat{\ell}\}) = \mathcal{P}_* \equiv \mathcal{F}_n(\{\ell_1\}) \cap \mathcal{F}_n(\{\ell_2\}) \cap \dots \cap \mathcal{F}_n(\{\ell_R\}).$$

*Equivalently,  $\{\hat{\ell}\} \models \mathcal{B}_*p$  only if  $\{\ell_i\} \models \mathcal{B}_np$  for all  $i$ .*

# Our $t$ -defended random-search EHPO

**Definition 7.** Suppose that we are given a naive EHPO procedure  $(\{H\}, \mathcal{F}_n)$ , in which  $H$  is random search and is the only HPO in our EHPO, and  $\mathcal{F}_n$  is a “naive” belief function associated with a naive reasoner  $\mathcal{B}_n$ . For any  $K, R \in \mathbb{N}$ , we define the “ $(K, R)$ -defended” belief function  $\mathcal{F}_*$  for a skeptical reasoner  $\mathcal{B}_*$  as the following conclusion-drawing procedure. First,  $\mathcal{F}_*$  only makes conclusion set  $\mathcal{P}_*$  from a single log  $\hat{\ell}$  with  $K * R$  trials; otherwise, it concludes nothing, outputting  $\emptyset$ . Second,  $\mathcal{F}_*$  splits the single  $\hat{\ell}$  into  $R$  logs  $\ell_1, \ell_2, \dots, \ell_R$ , each containing  $K$  independent-random-search trials.<sup>8</sup> Finally,  $\mathcal{F}_*$  outputs the intersection of what the naive reasoner would have output on each log  $\ell_i$ ,

$$\mathcal{F}_*(\{\hat{\ell}\}) = \mathcal{P}_* \equiv \mathcal{F}_n(\{\ell_1\}) \cap \mathcal{F}_n(\{\ell_2\}) \cap \dots \cap \mathcal{F}_n(\{\ell_R\}).$$

Equivalently,  $\{\hat{\ell}\} \models \mathcal{B}_*p$  only if  $\{\ell_i\} \models \mathcal{B}_np$  for all  $i$ .



# Our $t$ -defended random-search EHPO

**Definition 7.** Suppose that we are given a naive EHPO procedure  $(\{H\}, \mathcal{F}_n)$ , in which  $H$  is random search and is the only HPO in our EHPO, and  $\mathcal{F}_n$  is a “naive” belief function associated with a naive reasoner  $\mathcal{B}_n$ . For any  $K, R \in \mathbb{N}$ , we define the “ $(K, R)$ -defended” belief function  $\mathcal{F}_*$  for a skeptical reasoner  $\mathcal{B}_*$  as the following conclusion-drawing procedure. First,  $\mathcal{F}_*$  only makes conclusion set  $\mathcal{P}_*$  from a single log  $\hat{\ell}$  with  $K * R$  trials; otherwise, it concludes nothing, outputting  $\emptyset$ . Second,  $\mathcal{F}_*$  splits the single  $\hat{\ell}$  into  $R$  logs  $\ell_1, \ell_2, \dots, \ell_R$ , each containing  $K$  independent-random-search trials.<sup>8</sup> Finally,  $\mathcal{F}_*$  outputs the intersection of what the naive reasoner would have output on each log  $\ell_i$ ,

$$\mathcal{F}_*(\{\hat{\ell}\}) = \mathcal{P}_* \equiv \mathcal{F}_n(\{\ell_1\}) \cap \mathcal{F}_n(\{\ell_2\}) \cap \dots \cap \mathcal{F}_n(\{\ell_R\}).$$

Equivalently,  $\{\hat{\ell}\} \models \mathcal{B}_*p$  only if  $\{\ell_i\} \models \mathcal{B}_np$  for all  $i$ .

# Our $t$ -defended random-search EHPO

**Definition 7.** Suppose that we are given a naive EHPO procedure  $(\{H\}, \mathcal{F}_n)$ , in which  $H$  is random search and is the only HPO in our EHPO, and  $\mathcal{F}_n$  is a “naive” belief function associated with a naive reasoner  $\mathcal{B}_n$ . For any  $K, R \in \mathbb{N}$ , we define the “ $(K, R)$ -defended” belief function  $\mathcal{F}_*$  for a skeptical reasoner  $\mathcal{B}_*$  as the following conclusion-drawing procedure. First,  $\mathcal{F}_*$  only makes conclusion set  $\mathcal{P}_*$  from a single log  $\hat{\ell}$  with  $K * R$  trials; otherwise, it concludes nothing, outputting  $\emptyset$ . Second,  $\mathcal{F}_*$  splits the single  $\hat{\ell}$  into  $R$  logs  $\ell_1, \ell_2, \dots, \ell_R$ , each containing  $K$  independent-random-search trials.<sup>8</sup> Finally,  $\mathcal{F}_*$  outputs the intersection of what the naive reasoner would have output on each log  $\ell_i$ ,

$$\mathcal{F}_*(\{\hat{\ell}\}) = \mathcal{P}_* \equiv \mathcal{F}_n(\{\ell_1\}) \cap \mathcal{F}_n(\{\ell_2\}) \cap \dots \cap \mathcal{F}_n(\{\ell_R\}).$$

Equivalently,  $\{\hat{\ell}\} \models \mathcal{B}_*p$  only if  $\{\ell_i\} \models \mathcal{B}_np$  for all  $i$ .

# Our $t$ -defended random-search EHPO

**Informally**, to draw a conclusion using this EHPO,  $\mathcal{B}_*$  splits a random-search-trial log of size  $K * R$  into  $R$  groups of  $K$ -trial logs, passing each  $K$ -trial log to one of an ensemble of  $R$  naive reasoners  $\mathcal{B}_n$ .  $\mathcal{B}_*$  only concludes  $p$  if all  $R$  naive reasoners unanimously agree on  $p$ . We can guarantee this EHPO to be  $t$ -non-deceptive by assuming a bound on how much the hyper-HPs can affect the HPs.

# Our $t$ -defended random-search EHPO

**Informally**, to draw a conclusion using this EHPO,  $\mathcal{B}_*$  splits a random-search-trial log of size  $K * R$  into  $R$  groups of  $K$ -trial logs, passing each  $K$ -trial log to one of an ensemble of  $R$  naive reasoners  $\mathcal{B}_n$ .  $\mathcal{B}_*$  only concludes  $p$  if all  $R$  naive reasoners unanimously agree on  $p$ . We can guarantee this EHPO to be  $t$ -non-deceptive by assuming a bound on how much the hyper-HPs can affect the HPs.



# Our $t$ -defended random-search EHPO

**Informally**, to draw a conclusion using this EHPO,  $\mathcal{B}_*$  splits a random-search-trial log of size  $K * R$  into  $R$  groups of  $K$ -trial logs, passing each  $K$ -trial log to one of an ensemble of  $R$  naive reasoners  $\mathcal{B}_n$ .  $\mathcal{B}_*$  only concludes  $p$  if all  $R$  naive reasoners unanimously agree on  $p$ . We can guarantee this EHPO to be  $t$ -non-deceptive by assuming a bound on how much the hyper-HPs can affect the HPs.

**Theorem 1.** *Suppose that the set of allowable hyper-HPs  $\mathcal{C}$  of  $H$  is constrained, such that any two allowable random-search distributions  $\mu$  and  $\nu$  have Renyi- $\infty$ -divergence at most a constant, i.e.  $D_\infty(\mu\|\nu) \leq \gamma$ . The  $(K, R)$ -defended random-search EHPO of Definition 7 is guaranteed to be  $t$ -non-deceptive if we set  $R \geq \sqrt{t \exp(\gamma K)/K} = O(\sqrt{t})$ .*

# Our $t$ -defended EHPO can make conclusions using fewer than $t$ resources

We describe a defended variation of **random search**

Our defended-random-search EHPO is defended against deception **for time budget  $t$** , but it is able to draw conclusions **using up compute budgets that are  $O(\sqrt{t})$**

In short, our algorithm can make conclusions **using much fewer resources** than the total compute budget for which it is defended against deception (i.e.,  $t$ -defended in runtime on the order of  $\sqrt{t}$ )

# Our $t$ -defended random-search EHPO

---

**Algorithm 1** Defense with Random Search

---

**Require:** Set of  $K * R$  random-search logs  $\{\mathcal{L}_i\}_{i=1}^{KR}$ ,  
defense subsampling budget  $M$ ,  
criterion constant  $\delta$ , subsample size  $\kappa$ .

- 1: **for**  $m = 1, \dots, M$  **do**
  - 2:   Subsample  $\kappa$  logs:  $\{\mathcal{L}_i\}_{i=1}^{\kappa} \sim \{\mathcal{L}_i\}_{i=1}^{KR}$ .
  - 3:   Obtain conclusions  $\{\mathcal{P}_i\}_{i=1}^{\kappa}$  from  $\{\mathcal{L}_i\}_{i=1}^{\kappa}$ .
  - 4:   Obtain output conclusion for  $m$ :  
       $\mathcal{P}^{(m)} \leftarrow \text{Majority}(\{\mathcal{P}_i\}_{i=1}^{\kappa})$
  - 5: **end for**
  - 6: **if**  $\exists p$  s.t.  $\geq (1 - \delta)M$  of  $\{\mathcal{P}^{(m)}\}_{i=1}^M$  **conclude**  $p$   
   **then**
  - 7:   Conclude  $p$ .
  - 8: **else**
  - 9:   Conclude nothing.
  - 10: **end if**
- 

$K * R$ -size super log

---

# Our $t$ -defended random-search EHPO

---

**Algorithm 1** Defense with Random Search

---

**Require:** Set of  $K * R$  random-search logs  $\{\mathcal{L}_i\}_{i=1}^{KR}$ ,  
defense subsampling budget  $M$ ,  
criterion constant  $\delta$ , subsample size  $\kappa$ .

```
1: for  $m = 1, \dots, M$  do
2:   Subsample  $\kappa$  logs:  $\{\mathcal{L}_i\}_{i=1}^{\kappa} \sim \{\mathcal{L}_i\}_{i=1}^{KR}$ .
3:   Obtain conclusions  $\{\mathcal{P}_i\}_{i=1}^{\kappa}$  from  $\{\mathcal{L}_i\}_{i=1}^{\kappa}$ .
4:   Obtain output conclusion for  $m$ :
        $\mathcal{P}^{(m)} \leftarrow \text{Majority}(\{\mathcal{P}_i\}_{i=1}^{\kappa})$ 
5: end for
6: if  $\exists p$  s.t.  $\geq (1 - \delta)M$  of  $\{\mathcal{P}^{(m)}\}_{i=1}^M$  conclude  $p$ 
   then
7:   Conclude  $p$ .
8: else
9:   Conclude nothing.
10: end if
```

---



For  $M$  iterations,  
we subsample  $\kappa$  logs,  
and pass them to  
 $\kappa$  naive reasoners  $B_n$

Each iteration  $m$  concludes  
the **majority conclusion**  
of the  $\kappa$ -sized  $B_n$  ensemble

# Our $t$ -defended random-search EHPO

---

**Algorithm 1** Defense with Random Search

---

**Require:** Set of  $K * R$  random-search logs  $\{\mathcal{L}_i\}_{i=1}^{KR}$ ,  
defense subsampling budget  $M$ ,  
criterion constant  $\delta$ , subsample size  $\kappa$ .

```
1: for  $m = 1, \dots, M$  do
2:   Subsample  $\kappa$  logs:  $\{\mathcal{L}_i\}_{i=1}^{\kappa} \sim \{\mathcal{L}_i\}_{i=1}^{KR}$ .
3:   Obtain conclusions  $\{\mathcal{P}_i\}_{i=1}^{\kappa}$  from  $\{\mathcal{L}_i\}_{i=1}^{\kappa}$ .
4:   Obtain output conclusion for  $m$ :
        $\mathcal{P}^{(m)} \leftarrow \text{Majority}(\{\mathcal{P}_i\}_{i=1}^{\kappa})$ 
5: end for
6: if  $\exists p$  s.t.  $\geq (1 - \delta)M$  of  $\{\mathcal{P}^{(m)}\}_{i=1}^M$  conclude  $p$ 
   then
7:   Conclude  $p$ .
8: else
9:   Conclude nothing.
10: end if
```

---

If we meet a certain number  
of  $M$ -majority conclusions,

we conclude  $p$

# Our $t$ -defended random-search EHPO

---

## Algorithm 1 Defense with Random Search

---

**Require:** Set of  $K * R$  random-search logs  $\{\mathcal{L}_i\}_{i=1}^{KR}$ ,  
 defense subsampling budget  $M$ ,  
 criterion constant  $\delta$ , subsample size  $\kappa$ .

- 1: **for**  $m = 1, \dots, M$  **do**
- 2:   Subsample  $\kappa$  logs:  $\{\mathcal{L}_i\}_{i=1}^{\kappa} \sim \{\mathcal{L}_i\}_{i=1}^{KR}$ .
- 3:   Obtain conclusions  $\{\mathcal{P}_i\}_{i=1}^{\kappa}$  from  $\{\mathcal{L}_i\}_{i=1}^{\kappa}$ .
- 4:   Obtain output conclusion for  $m$ :  
        $\mathcal{P}^{(m)} \leftarrow \text{Majority}(\{\mathcal{P}_i\}_{i=1}^{\kappa})$
- 5: **end for**
- 6: **if**  $\exists p$  s.t.  $\geq (1 - \delta)M$  of  $\{\mathcal{P}^{(m)}\}_{i=1}^M$  **conclude**  $p$   
    **then**
- 7:   Conclude  $p$ .
- 8: **else**
- 9:   Conclude nothing.
- 10: **end if**

---


Table 1: Results from repeating our Section 2 experiment, using Algorithm 1 instead of grid search.  $p$  = “Non-adaptive optimizers (SGD and HB) perform better than the adaptive optimizer Adam”.

	$p$	$\neg p$	$1 - \delta$	Conclude
SGD vs. Adam	0.213	0.788	0.75	$\neg p$
			0.8	Nothing
			0.9	Nothing
HB vs. Adam	0.168	0.832	0.75	$\neg p$
			0.8	$\neg p$
			0.9	Nothing

# We empirically validate our suggested $t$ -non-deceptive EHPO



$p$





$\neg p$

	$p$	$\neg p$	$1 - \delta$	Conclude
SGD vs. Adam	0.213	0.788	0.75	$\neg p$
			0.8	Nothing
			0.9	Nothing
HB vs. Adam	0.168	0.832	0.75	$\neg p$
			0.8	$\neg p$
			0.9	Nothing



# We empirically validate our suggested $t$ -non-deceptive EHPO





	$p$	$\neg p$	$1 - \delta$	Conclude
SGD vs. Adam	0.213	0.788	0.75	$\neg p$
			0.8	Nothing
			0.9	Nothing
HB vs. Adam	0.168	0.832	0.75	$\neg p$
			0.8	$\neg p$
			0.9	Nothing

Our belief / skepticism



# We empirically validate our suggested $t$ -non-deceptive EHPO



	$p$	$\neg p$	$1 - \delta$	Conclude
SGD vs. Adam	0.213	0.788	0.75	$\neg p$
			0.8	Nothing
			0.9	Nothing
HB vs. Adam	0.168	0.832	0.75	$\neg p$
			0.8	$\neg p$
			0.9	Nothing

Our belief / skepticism

# Takeaways for Robust ML

It is possible to construct  $t$ -defended EHPO, such as our  $t$ -defended random search

# Takeaways for Robust ML


It is possible to construct  $t$ -defended EHPO, such as our  $t$ -defended random search

Any defense will depend on assumptions concerning how we configure underlying HPO, so researchers must be explicit about their configuration choices and their notion of belief/skepticism


# Takeaways for Robust ML

It is possible to construct  $t$ -defended EHPO, such as our  $t$ -defended random search

Any defense will depend on assumptions concerning how we configure underlying HPO, so researchers must be explicit about their configuration choices and their notion of **belief/skepticism**



$p$



$\neg p$

# Takeaways for Robust ML

It is possible to construct  $t$ -defended EHPO, such as our  $t$ -defended random search

Any defense will depend on assumptions concerning how we configure underlying HPO, so researchers must be explicit about their configuration choices and their notion of belief/skepticism

Avoiding deception is just as important as ensuring reproducibility, as we want to ensure results are both replicable and correct


# Future work

**Plug-and-play** defended EHPO for a specific class of ML algorithms

**Leverages the logic** to suggest a specific, defended grid-search-based hyper-HP selection algorithm (rather than a general possibility proof)

**Takes an adversarial approach** to defining our notion of skepticism

**Provides a clear recipe** for developing  $t$ -defended EHPO for different ML domains



# Cornell Bowers C-IS

College of Computing  
and Information Science

**Thank you!**